

Homework 2 Solutions

Ian Cranfield [REDACTED]

4/21/2020 by 5PM

```
library(dynlm)
library(pastecs)
library(openxlsx)
library(tidyverse)
library(data.table)
library(scales)
library(TSstudio)
library("readxl")
library(ggplot2)
require(cowplot)
library(psych)
library(corrplot)
library(gtable)
library(timeDate)
library(PerformanceAnalytics)
library(fpp2)
library(seasonal)
```

Problem 4.3:


```
#### Problem 4.3 ####
Quarter_house <- read_excel("Chapter4_exercises_data.xls", sheet =2)
setnames(Quarter_house, c('Quarter', 'P', 'R'))
Quarter_house$Quarter <- as.Date( as.yearqtr(Quarter_house$Quarter, format = "%YQ%q" ))
#Price <- ts(Quarter_house$P, start = 1980.25, frequency = 4)
#Rate <- ts(Quarter_house$R, start = 1980.25, frequency = 4)

glimpse(Quarter_house)

## Rows: 126
## Columns: 3
## $ Quarter <date> 1980-04-01, 1980-07-01, 1980-10-01, 1981-01-01, 1981-04-01...
## $ P      <dbl> 42.49871, 43.27341, 43.78586, 44.35695, 45.07844, 45.55761,...
## $ R      <dbl> 14.43000, 12.65000, 14.26333, 15.14333, 16.22667, 17.42333,...

acf_price <- acf(Quarter_house$P, plot = T, main = "ACF of House Prices")
```

ACF of House Prices




```
acf_price
```

```
##
## Autocorrelations of series 'Quarter_house$P', by lag
##
##      0      1      2      3      4      5      6      7      8      9      10     11     12
## 1.000 0.988 0.974 0.959 0.941 0.921 0.898 0.876 0.852 0.826 0.799 0.772 0.743
##      13     14     15     16     17     18     19     20     21
## 0.713 0.682 0.651 0.619 0.585 0.552 0.519 0.486 0.453
```

```
pacf_price <- pacf(Quarter_house$P, plot= T, main = "PACF of House Prices")
```


PACF of House Prices



```
pacf_price
```

```
##
## Partial autocorrelations of series 'Quarter_house$P', by lag
##
##      1      2      3      4      5      6      7      8      9     10     11
## 0.988 -0.098 -0.033 -0.100 -0.126 -0.067  0.000 -0.052 -0.072 -0.020  0.006
##      12     13     14     15     16     17     18     19     20     21
## -0.068 -0.087 -0.019  0.006 -0.045 -0.055 -0.014  0.023  0.002 -0.030
acf_rate <- acf(Quarter_house$R, plot = TRUE, main = "ACF of Interest Rates")
```

ACF of Interest Rates




```
acf_rate
```

```
##
## Autocorrelations of series 'Quarter_house$R', by lag
##
##      0      1      2      3      4      5      6      7      8      9      10     11     12
## 1.000 0.966 0.934 0.895 0.850 0.802 0.757 0.716 0.681 0.648 0.615 0.590 0.568
## 13     14     15     16     17     18     19     20     21
## 0.547 0.519 0.491 0.465 0.437 0.407 0.384 0.367 0.351
```

```
pacf_rate <- pacf(Quarter_house$R, plot = TRUE, main = "PACF of Interest Rates")
```

PACF of Interest Rates




```
pacf_rate
```

```
##
## Partial autocorrelations of series 'Quarter_house$R', by lag
##
##      1      2      3      4      5      6      7      8      9     10     11
## 0.966 0.024 -0.135 -0.122 -0.052 0.033 0.062 0.057 0.009 -0.069 0.084
##      12     13     14     15     16     17     18     19     20     21
## 0.037 -0.013 -0.124 -0.060 0.034 0.014 -0.035 0.093 0.060 0.003
```

```
Quarter_house <- Quarter_house %>%
  mutate( P_growth = log(P)-log(lag(P)))
```

```
acf_price_growth <- acf(Quarter_house$P_growth, plot = TRUE, na.action = na.omit, main = "ACF of House P Growth")
```

ACF of House Price Growth



```
acf_price_growth
```

```
##
```

```
## Autocorrelations of series 'Quarter_house$P_growth', by lag
```

```
##
```

```
##      0      1      2      3      4      5      6      7      8      9     10
```


```
## 1.000 0.628 0.294 0.498 0.705 0.372 0.121 0.325 0.520 0.250 -0.007
```

```
##    11    12    13    14    15    16    17    18    19    20
```

```
## 0.139 0.283 0.014 -0.224 -0.078 0.098 -0.111 -0.309 -0.156 0.011
```

```
pacf_price_growth <- pacf(Quarter_house$P_growth, plot = TRUE, na.action = na.omit, main = "PACF of House Price Growth")
```

PACF of House Price Growth




```
pacf_price_growth
```

```
##
## Partial autocorrelations of series 'Quarter_house$P_growth', by lag
##
##      1      2      3      4      5      6      7      8      9     10     11
## 0.628 -0.165 0.655 0.195 -0.427 0.111 0.024 0.165 -0.106 -0.203 -0.025
##      12     13     14     15     16     17     18     19     20
## -0.103 -0.130 -0.027 -0.071 0.126 -0.090 0.015 0.020 -0.039
```

```
acf_rate_change <- acf(diff(Quarter_house$R), na.action = na.omit, main = "ACF of Interest Rate Changes")
```

ACF of Interest Rate Changes



```
acf_rate_change
```

```
##
```

```
## Autocorrelations of series 'diff(Quarter_house$R)', by lag
```

```
##
```

```
##      0      1      2      3      4      5      6      7      8      9     10
```


```
## 1.000 0.220 0.035 -0.005 -0.126 -0.268 -0.232 -0.175 -0.006 0.090 0.008
```

```
##     11     12     13     14     15     16     17     18     19     20
```

```
## 0.156 0.063 0.011 0.020 -0.014 -0.045 -0.032 -0.183 -0.104 -0.006
```

```
pacf_rate_change <- pacf(diff(Quarter_house$R), na.action = na.omit, main = "PACF of Interest Rate Changes")
```


PACF of Interest Rate Changes



```
pacf_rate_change
```

```
##
## Partial autocorrelations of series 'diff(Quarter_house$R)', by lag
##
##      1      2      3      4      5      6      7      8      9     10     11
## 0.220 -0.014 -0.010 -0.129 -0.227 -0.146 -0.118  0.036  0.044 -0.113  0.076
##      12     13     14     15     16     17     18     19     20
## -0.077 -0.016  0.032 -0.004  0.007 -0.019 -0.187 -0.048 -0.014
```

We observed the ACF and PACF of Quarterly House Prices, Interest Rates, House Price Growth, and Interest Rate Changes.

The ACF of House Price and Interest Rates are significantly correlated but slowly decreases to 0. But, the ACF of House Price Growth are significant correlated for the first 12 lags.

The PACF of House Price and Interest Rates remain stationary with no significant correlations meaning the past Quarter does not correlate with the present Quarter. The PACF of House Price Growth are significantly correlated for the first 5 lags. The PACF of Interest Rate Changes are significantly correlated at lags 5, and 18 where it is less than -0.2.

Problem 4.4:

```
## 4.4
# Annual data from 4.1
Annual_house <- read_excel("Chapter4_exercises_data.xls", sheet =1)
setnames(Annual_house, c("Year", "P", "R"))
```

```
ts_P_annual <- ts(Annual_house$P, start = 1975, frequency = 1)
```

```
Ann_P_growth <- diff(log(ts_P_annual))
```

```
reg44iA <- dynlm(Ann_P_growth ~ L(Ann_P_growth,1))
summary(reg44iA)
```

```
##
## Time series regression with "ts" data:
## Start = 1977, End = 2011
##
## Call:
## dynlm(formula = Ann_P_growth ~ L(Ann_P_growth, 1))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.086311 -0.013786  0.007275  0.017808  0.050339
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.001363   0.006727   0.203   0.841
## L(Ann_P_growth, 1) 0.889768   0.099457   8.946 2.44e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.02913 on 33 degrees of freedom
## Multiple R-squared:  0.7081, Adjusted R-squared:  0.6992
## F-statistic: 80.03 on 1 and 33 DF,  p-value: 2.437e-10
```

```
reg44iiA <- dynlm(Ann_P_growth ~ L(Ann_P_growth,1:2))
summary(reg44iiA)
```

```
##
## Time series regression with "ts" data:
## Start = 1978, End = 2011
##
## Call:
## dynlm(formula = Ann_P_growth ~ L(Ann_P_growth, 1:2))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.075902 -0.008761  0.006968  0.015660  0.031019
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.006284   0.006204   1.013  0.31894
## L(Ann_P_growth, 1:2)1  1.269847   0.157396   8.068 4.13e-09 ***
## L(Ann_P_growth, 1:2)2 -0.485851   0.161879  -3.001  0.00527 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.02556 on 31 degrees of freedom
## Multiple R-squared:  0.7791, Adjusted R-squared:  0.7649
## F-statistic: 54.68 on 2 and 31 DF,  p-value: 6.823e-11
```

```
reg44iiiA <- dynlm(Ann_P_growth ~ L(Ann_P_growth,1:3))
summary(reg44iiiA)
```

```
##
## Time series regression with "ts" data:
## Start = 1979, End = 2011
##
## Call:
## dynlm(formula = Ann_P_growth ~ L(Ann_P_growth, 1:3))
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-0.070274	-0.013101	0.005806	0.017802	0.031248

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.002655	0.008284	0.320	0.7509
L(Ann_P_growth, 1:3)1	1.339177	0.239772	5.585	5e-06 ***
L(Ann_P_growth, 1:3)2	-0.679351	0.388512	-1.749	0.0909 .
L(Ann_P_growth, 1:3)3	0.174552	0.277536	0.629	0.5343

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.02588 on 29 degrees of freedom
## Multiple R-squared:  0.7674, Adjusted R-squared:  0.7433
## F-statistic: 31.89 on 3 and 29 DF,  p-value: 2.551e-09
```

```
reg44ivA <- dynlm(Ann_P_growth ~ L(Ann_P_growth,1:4))
summary(reg44ivA)
```

```
##
## Time series regression with "ts" data:
## Start = 1980, End = 2011
##
## Call:
## dynlm(formula = Ann_P_growth ~ L(Ann_P_growth, 1:4))
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-0.065867	-0.010847	0.005646	0.017617	0.035556

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.006745	0.011337	0.595	0.5568
L(Ann_P_growth, 1:4)1	1.317152	0.251458	5.238	1.61e-05 ***
L(Ann_P_growth, 1:4)2	-0.749136	0.420709	-1.781	0.0862 .
L(Ann_P_growth, 1:4)3	0.353245	0.433174	0.815	0.4219
L(Ann_P_growth, 1:4)4	-0.164668	0.302110	-0.545	0.5902

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.02666 on 27 degrees of freedom
## Multiple R-squared:  0.7489, Adjusted R-squared:  0.7117
## F-statistic: 20.13 on 4 and 27 DF,  p-value: 8.783e-08
```

```
BIC(reg44iA, reg44iiiA, reg44iiiiA, reg44ivA)
```

```
## Warning in BIC.default(reg44iA, reg44iiiA, reg44iiiiA, reg44ivA): models are not
## all fitted to the same number of observations
```

```
##           df          BIC
## reg44iA    3 -139.5969
## reg44iiiA  4 -141.8926
## reg44iiiiA 5 -134.3214
## reg44ivA   6 -125.8103
```

```
AIC(reg44iA, reg44iiiA, reg44iiiiA, reg44ivA)
```

```
## Warning in AIC.default(reg44iA, reg44iiiA, reg44iiiiA, reg44ivA): models are not
## all fitted to the same number of observations
```

```
##           df          AIC
## reg44iA    3 -144.2630
## reg44iiiA  4 -147.9981
## reg44iiiiA 5 -141.8039
## reg44ivA   6 -134.6047
```

```
#Choose model 2
```

The second model displays the lowest values for BIC and AIC so Model 2 is preferred.

Quarterly Regression

```
## Quarter Regression
```

```
Quarterly <- read_excel("Chapter4_exercises_data.xls", sheet =2)
```

```
ts_P_Quarter <- ts(Quarterly$P, start = 1980.25, frequency =4)
```

```
Q_P_growth <- diff(log(ts_P_Quarter))
```

```
reg44iq <- dynlm(Q_P_growth ~ L(Q_P_growth,1))
summary(reg44iq)
```

```
##
## Time series regression with "ts" data:
## Start = 1980(4), End = 2011(3)
##
## Call:
## dynlm(formula = Q_P_growth ~ L(Q_P_growth, 1))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.048408 -0.006184  0.000542  0.006037  0.036627
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.002937   0.001257   2.336   0.0211 *
## L(Q_P_growth, 1) 0.632166   0.070382   8.982 4.03e-15 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 0.0123 on 122 degrees of freedom
## Multiple R-squared: 0.3981, Adjusted R-squared: 0.3931
## F-statistic: 80.68 on 1 and 122 DF, p-value: 4.034e-15

reg44iiq <- dynlm(Q_P_growth ~ L(Q_P_growth,1:2))
summary(reg44iiq)

##
## Time series regression with "ts" data:
## Start = 1981(1), End = 2011(3)
##
## Call:
## dynlm(formula = Q_P_growth ~ L(Q_P_growth, 1:2))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.048606 -0.005667  0.000845  0.007060  0.031736
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.003500   0.001282   2.731  0.00728 **
## L(Q_P_growth, 1:2)1  0.743849   0.090491   8.220 2.74e-13 ***
## L(Q_P_growth, 1:2)2 -0.174611   0.090360  -1.932  0.05567 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.01222 on 120 degrees of freedom
## Multiple R-squared: 0.4162, Adjusted R-squared: 0.4065
## F-statistic: 42.78 on 2 and 120 DF, p-value: 9.454e-15

reg44iiiq <- dynlm(Q_P_growth ~ L(Q_P_growth,1:3))
summary(reg44iiiq)

##
## Time series regression with "ts" data:
## Start = 1981(2), End = 2011(3)
##
## Call:
## dynlm(formula = Q_P_growth ~ L(Q_P_growth, 1:3))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.033598 -0.003980  0.000351  0.004362  0.023351
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.0003837   0.0009288   0.413   0.68
## L(Q_P_growth, 1:3)1  0.9018739   0.0641574  14.057 < 2e-16 ***
## L(Q_P_growth, 1:3)2 -0.7623828   0.0808082  -9.434 4.38e-16 ***
## L(Q_P_growth, 1:3)3  0.7633709   0.0664357  11.490 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.008459 on 118 degrees of freedom
```

```
## Multiple R-squared:  0.7245, Adjusted R-squared:  0.7175
## F-statistic: 103.4 on 3 and 118 DF,  p-value: < 2.2e-16
```

```
reg44ivq <- dynlm(Q_P_growth ~ L(Q_P_growth,1:4))
summary(reg44ivq)
```

```
##
## Time series regression with "ts" data:
## Start = 1981(3), End = 2011(3)
##
## Call:
## dynlm(formula = Q_P_growth ~ L(Q_P_growth, 1:4))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.034072 -0.003491  0.000696  0.003942  0.021627
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    9.883e-05  9.383e-04   0.105   0.9163
## L(Q_P_growth, 1:4)1  7.815e-01  9.217e-02   8.479 8.39e-14 ***
## L(Q_P_growth, 1:4)2 -6.370e-01  1.061e-01  -6.001 2.29e-08 ***
## L(Q_P_growth, 1:4)3  6.087e-01  1.081e-01   5.630 1.28e-07 ***
## L(Q_P_growth, 1:4)4  1.745e-01  9.618e-02   1.814  0.0723 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.008413 on 116 degrees of freedom
## Multiple R-squared:  0.7316, Adjusted R-squared:  0.7223
## F-statistic: 79.04 on 4 and 116 DF,  p-value: < 2.2e-16
```

```
BIC(reg44iq, reg44iiq, reg44iiiq, reg44ivq)
```

```
## Warning in BIC.default(reg44iq, reg44iiq, reg44iiiq, reg44ivq): models are not
## all fitted to the same number of observations
```

```
##          df      BIC
## reg44iq    3 -726.3076
## reg44iiq    4 -718.3639
## reg44iiiq   5 -798.3111
## reg44ivq    6 -789.2122
```

```
AIC(reg44iq, reg44iiq, reg44iiiq, reg44ivq)
```

```
## Warning in AIC.default(reg44iq, reg44iiq, reg44iiiq, reg44ivq): models are not
## all fitted to the same number of observations
```

```
##          df      AIC
## reg44iq    3 -734.7684
## reg44iiq    4 -729.6127
## reg44iiiq   5 -812.3312
## reg44ivq    6 -805.9869
```

```
# Choose model up to 3, reg44iiiq
```

Model 3 has the lowest BIC and AIC, so we would prefer using Model 3 that has up to 3 lags. All of the 3 lags are statistically significant as well. Makes more sense because it is quarterly data and the lags have more importance. I will be using Model 3 for the Recursive and Rolling Scheme.


Recursive Scheme*

```
recursive <- list()
endwith = 2000.25
loop_length <- NROW(Q_P_growth)-10*8.5

for (i in 1:loop_length) {
  recursive[[i]] <- dynlm(Q_P_growth ~ L(Q_P_growth,1:3), start=1980.5, end=endwith)$coefficients
  endwith <- 2000.25 + 0.25*i
}

recursive_df <- ts(data.frame(matrix(unlist(recursive), nrow=length(recursive), byrow = TRUE)),
  start= 2000.5, frequency = 4)

autoplot(recursive_df[,2], series = "Lag 1") +
  autolayer(recursive_df[,3], series = "Lag 2") +
  autolayer(recursive_df[,4], series = "Lag 3") +
  ggtitle("Recursive Linear Regression Coefficients") +
  ylab("Coefficients") +
  scale_color_manual(values = c("black", "blue", "red"), breaks = c("Lag 1", "Lag 2", "Lag 3"))
```



There is greater fluctuation in the beginning than it smooths out. There is another set of fluctuation around 2008, but this is because of the housing market crash in 2008. I used 2/3 of the data for the estimation sample so the prediction sample between 2005-2011 was the final 1/3.


Rolling Scheme*

```
rolling <- list()
startwith = 1980.25
loop_length <- NROW(Q_P_growth)-10*4

for (i in 1:loop_length) {
  rolling[[i]] <- dynlm(Q_P_growth ~ L(Q_P_growth,1:3), start=startwith, end=startwith+10)$coef
  startwith <- 1980.25 + 0.25*i
}

rolling_df <- ts(data.frame(matrix(unlist(rolling), nrow=length(rolling), byrow = TRUE)),
  start= 1980.5, frequency = 4)

autoplot(rolling_df[,2], series = "Lag 1") +
  autolayer(rolling_df[,3], series = "Lag 2") +
  autolayer(rolling_df[,4], series = "Lag 3") +
  ggtitle("Rolling Linear Regression Coefficients") +
  ylab("Coefficients") +
  scale_color_manual(values = c("black", "blue", "red"), breaks = c("Lag 1", "Lag 2", "Lag 3"))
```



The rolling scheme is pretty similar to the recursive scheme but the estimation window is fixed at 40 observations.

*Worked with Carmel Shek

Problem 4.8:

```
greenbook <- read_excel("Chapter4_exercises_data.xls", sheet=3)

setnames(greenbook, c('Date', 'actual_growth', 'forecasted_growth'))

greenbook$Date <- as.Date( as.yearqtr(greenbook$Date, format = "%Y:Q%q" ))

greenbook <- greenbook %>%
  mutate(
    forecast_errors = actual_growth - forecasted_growth
  )

errors = ts(greenbook$forecast_errors, start = 1969-01-01, frequency = 4)

regress_forecast_errors <- dynlm(errors ~ L(errors, 1:3) )

summary(regress_forecast_errors)
```

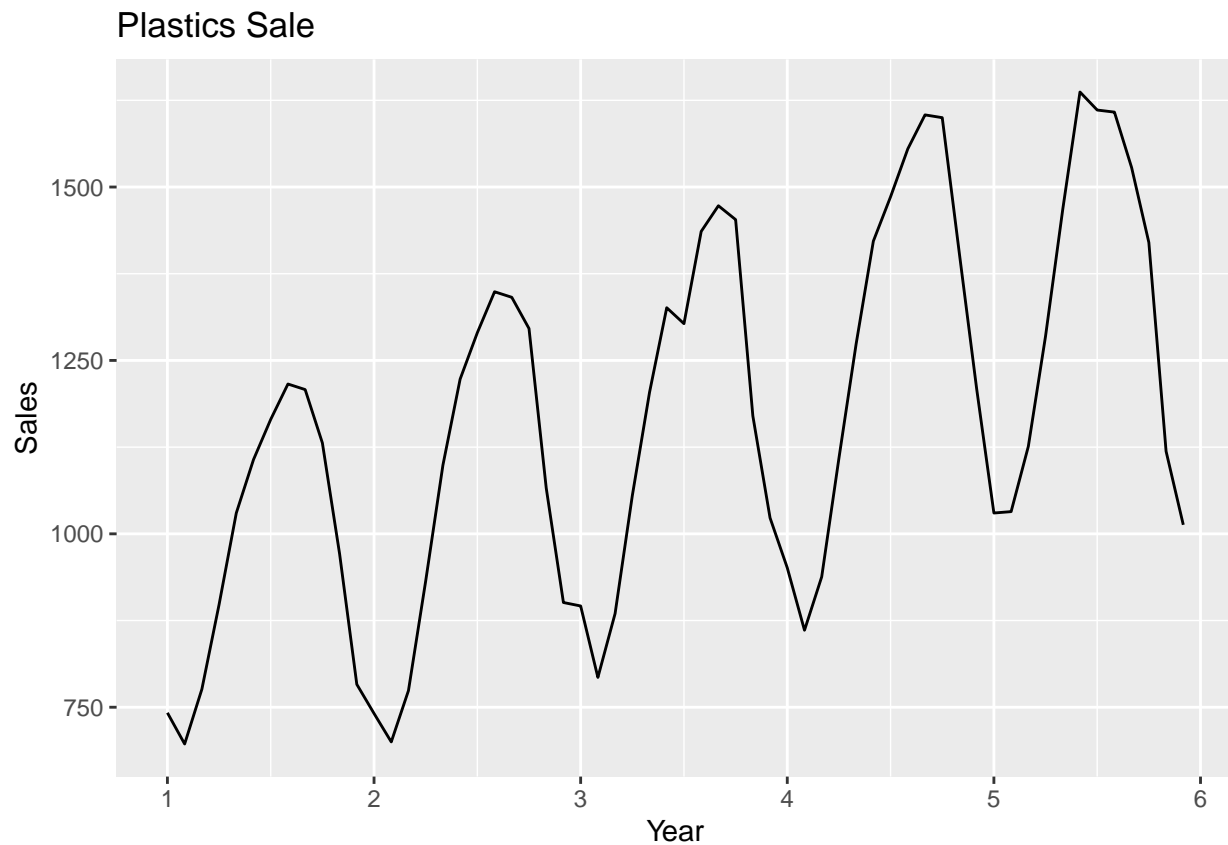
```
##
## Time series regression with "ts" data:
## Start = 1967(4), End = 2004(4)
##
## Call:
## dynlm(formula = errors ~ L(errors, 1:3))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.51624 -0.43181 -0.02828  0.35103  2.91904
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   0.060552   0.062702   0.966   0.336
## L(errors, 1:3)1 0.083250   0.083178   1.001   0.319
## L(errors, 1:3)2 0.054550   0.082762   0.659   0.511
## L(errors, 1:3)3 0.005634   0.082664   0.068   0.946
##
## Residual standard error: 0.7539 on 145 degrees of freedom
## Multiple R-squared:  0.0108, Adjusted R-squared:  -0.009666
## F-statistic: 0.5277 on 3 and 145 DF,  p-value: 0.6639
```

Ran a regression on the forecast errors with 3 lags. The regression that was run does a poor job to explain the forecast errors. The F-test is 0.5277 with a p-value of 66%. This shows that none of regressors can explain the forecast errors. So from this, we the forecast error is not predictable from its own past.

Problem 6.2 (Textbook C):

a. Plot the time series of sales of product A. Can you identify seasonal fluctuations and/or a trend-cycle?

```
autoplot(plastics) + ggtitle("Plastics Sale") + xlab("Year") + ylab("Sales")
```




I can identify an upward trend and seasonal patterns hitting peaks/ increased sales around August-September and dips/ decreased sales around January at the start of the year.

b. Use a classical multiplicative decomposition to calculate the trend-cycle and seasonal indices.

```
decompose_plastics <- decompose(plastics, "multiplicative")  
plot(decompose_plastics)
```

Decomposition of multiplicative time series



Seasonal Indices

```
decompose_plastics$seasonal
```

```
##           Jan           Feb           Mar           Apr           May           Jun           Jul
## 1 0.7670466 0.7103357 0.7765294 0.9103112 1.0447386 1.1570026 1.1636317
## 2 0.7670466 0.7103357 0.7765294 0.9103112 1.0447386 1.1570026 1.1636317
## 3 0.7670466 0.7103357 0.7765294 0.9103112 1.0447386 1.1570026 1.1636317
## 4 0.7670466 0.7103357 0.7765294 0.9103112 1.0447386 1.1570026 1.1636317
## 5 0.7670466 0.7103357 0.7765294 0.9103112 1.0447386 1.1570026 1.1636317
##           Aug           Sep           Oct           Nov           Dec
## 1 1.2252952 1.2313635 1.1887444 0.9919176 0.8330834
## 2 1.2252952 1.2313635 1.1887444 0.9919176 0.8330834
## 3 1.2252952 1.2313635 1.1887444 0.9919176 0.8330834
## 4 1.2252952 1.2313635 1.1887444 0.9919176 0.8330834
## 5 1.2252952 1.2313635 1.1887444 0.9919176 0.8330834
```

#Trend Indices

```
decompose_plastics$trend
```

```
##           Jan           Feb           Mar           Apr           May           Jun           Jul
## 1          NA          NA          NA          NA          NA          NA 976.9583
## 2 1000.4583 1011.2083 1022.2917 1034.7083 1045.5417 1054.4167 1065.7917
## 3 1117.3750 1121.5417 1130.6667 1142.7083 1153.5833 1163.0000 1170.3750
## 4 1208.7083 1221.2917 1231.7083 1243.2917 1259.1250 1276.5833 1287.6250
## 5 1374.7917 1382.2083 1381.2500 1370.5833 1351.2500 1331.2500          NA
##           Aug           Sep           Oct           Nov           Dec
## 1  977.0417  977.0833  978.4167  982.7083  990.4167
## 2 1076.1250 1084.6250 1094.3750 1103.8750 1112.5417
## 3 1175.5000 1180.5417 1185.0000 1190.1667 1197.0833
## 4 1298.0417 1313.0000 1328.1667 1343.5833 1360.6250
```

```
## 5      NA      NA      NA      NA      NA
```

I have displayed the decomposition of the Plastics Sales. And showed the indices for Seasonal and Trend.

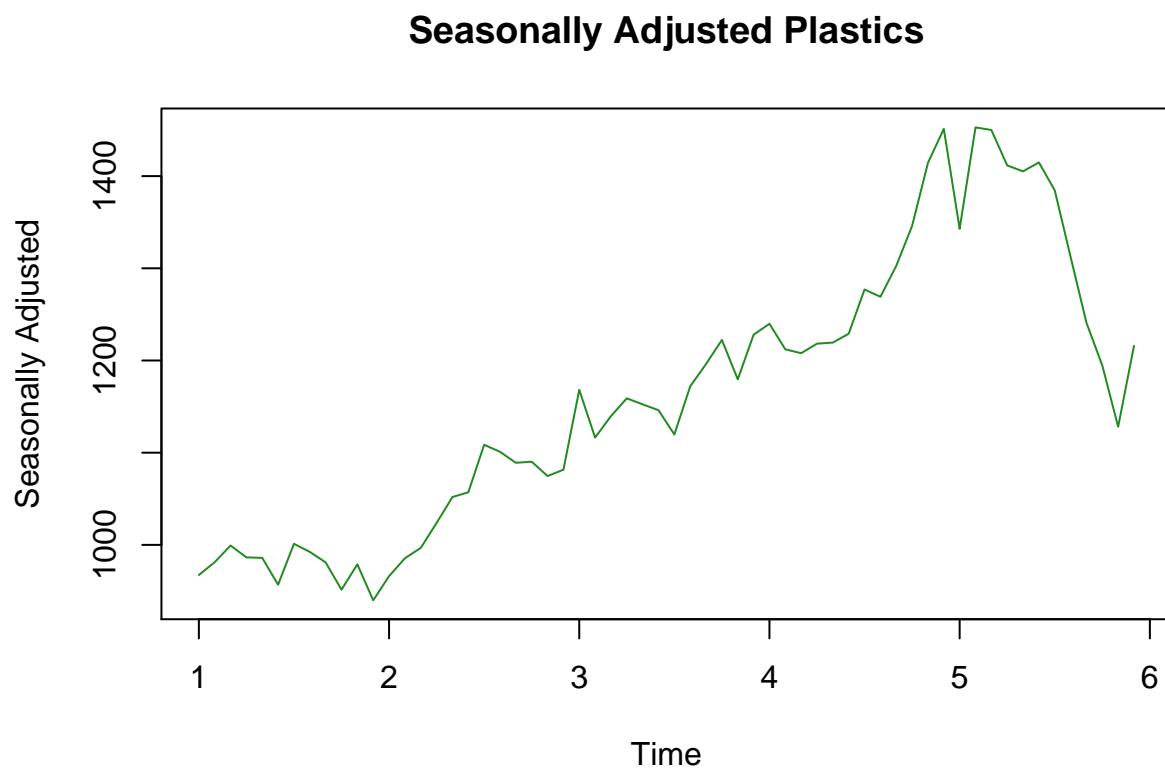
c. Do the results support the graphical interpretation from part a?

Yes, the results support the graphical interpretation in part A. The plot of the decomposed shows an upward trend and seasonal fluctuations

d. Compute and plot the seasonally adjusted data.

```
plastics_seas_adj <- seasadj(decompose_plastics)
```

```
plot(plastics_seas_adj, col = "forest green", ylab = "Seasonally Adjusted", xlab = "Time", main = "Seasonally Adjusted Plastics")
```



e. Change one observation to be an outlier (e.g., add 500 to one observation), and recompute the seasonally adjusted data. What is the effect of the outlier?

```
plastics2 <- plastics
```


```
plastics2[15] <- plastics2[15] + 500
```

```
decompose_plastics2 <- decompose(plastics2, "multiplicative")
```

```
plastics2_seas_adj <- seasadj(decompose_plastics2)
```

```
plot(plastics2_seas_adj, col = "black", ylab = "Seasonally Adjusted", xlab = "Time", main = "Seasonally Adjusted Plastics")
```

Seasonally Adjusted Plastics w/ outlier




Added 500 to a value to create an outlier. It created a spike where the outlier was inserted.

f. Does it make any difference if the outlier is near the end rather than in the middle of the time series?

```
plastics3 <- plastics  
  
plastics3[60] <- plastics[60] + 500  
  
decompose_plastics3 <- decompose(plastics3, "multiplicative")  
plastics3_seas_adj <- seasadj(decompose_plastics3)  
  
plot(plastics3_seas_adj, col = "black", ylab = "Seasonally Adjusted", xlab = "Time", main = "Seasonally  
Adjusted Plastics w/ outlier")  
lines(plastics_seas_adj, col="red", ylab="Seasonally Adjusted")
```

Seasonally Adjusted Plastics w/ outlier at the end



Added an outlier to the end of the time series, and it created a similar spike to the plot. We can see it follows the same pattern as the original data, but it has one spike at the end.

Problem 6.4 (Textbook C):

a. Write about 3–5 sentences describing the results of the decomposition. Pay particular attention to the scales of the graphs in making your interpretation.

When isolating the trend component from decomposition, it shows an overall upward trend indicating that the Australian Labour Force has increased over time. Around the early 1990s, it shows there were some dips to the labour force that can indicate a recession. The monthly breakdown shows that a few months had large variation and labour force increased around the months of March and December.


b. Is the recession of 1991/1992 visible in the estimated components?

Yes, the recessions in 1991/1992 are visible in the seasonally adjusted data.


Problem 6.5 (Textbook C):

a. Plot the data using `autoplot()`, `ggsubseriesplot()` and `ggseasonplot()` to look at the effect of the changing seasonality over time. What do you think is causing it to change so much?


```
autoplot(cangas)
```



```
ggsubseriesplot(cangas)
```



```
ggseasonplot(cangas)
```



In Winter months like January and December, gas production increased, but in summer months gas production decreased. Seen by analyzing the `ggsubseriesplot()`. This happens because the demand for gas probably increases during the Winter months because gas is probably used for heat. Gas production overall increased from 1960 to 2005, seen in the `ggseasonplot()`.


b. Do an STL decomposition of the data. You will need to choose `s.window` to allow for the changing shape of the seasonal component.

```
stl_cangas <- stl(cangas, s.window = 13, robust = TRUE)


autoplot(stl_cangas) + ggtitle("Monthly Canadian Gas Production", subtitle = "STL Decomposition")
```


Monthly Canadian Gas Production

STL Decomposition



```
autoplot(cangas, series = "Data") +
  autolayer(seasadj(stl_cangas), series = "Seasonally Adjusted") +
  autolayer(trendcycle(stl_cangas), series = "Trend" ) +
  ylab("Gas Production") +
  ggtitle("Monthly Canadian Gas Production", subtitle = "STL Decomposition") +
  scale_color_manual( values = c("grey", "blue", "red"), breaks = c("Data", "Seasonally Adjusted", "Trend"))
```




Shows an upward trend, and increase in seasonality in the 1980s.

c. Compare the results with those obtained using SEATS and X11. How are they different?

```
# X11
x11_cangas <- seas(cangas, x11 = "")
autoplot(x11_cangas) + ggtitle("Monthly Canadian Gas Production", subtitle = "X11 Decomposition")
```

Monthly Canadian Gas Production


X11 Decomposition



```
autoplot(cangas, series = "Data") +
  autolayer(seasadj(x11_cangas), series = "Seasonally Adjusted") +
  autolayer(trendcycle(x11_cangas), series = "Trend" ) +
  ylab("Gas Production") +
  ggtitle("Monthly Canadian Gas Production", subtitle = "X11 Decomposition") +
  scale_color_manual( values = c("grey", "blue", "red"), breaks = c("Data", "Seasonally Adjusted", "Trend"))
```

Monthly Canadian Gas Production

X11 Decomposition




```
#SEATS
```

```
seats_cangas <- seas(cangas)
autoplot(seats_cangas) + ggtitle("Monthly Canadian Gas Production", subtitle = "SEATS Decomposition")
```


Monthly Canadian Gas Production

SEATS Decomposition



```
autoplot(cangas, series = "Data") +
  autolayer(seasadj(seats_cangas), series = "Seasonally Adjusted") +
  autolayer(trendcycle(seats_cangas), series = "Trend") +
  ggtitle("Monthly Canadian Gas Production", subtitle = "SEATS Decomposition") +
  ylab("Gas Production") +
  scale_color_manual(values = c("gray", "blue", "red"), breaks = c("Data", "Seasonally Adjusted", "Trend"))
```

Monthly Canadian Gas Production SEATS Decomposition



The X11 and SEATS Decompositions are similar, and in the STL decomposition, the seasonally adjusted has more variance compared the X11 and SEATS Decompositions.


Problem 6.6 (Textbook C):

a. Use an STL decomposition to calculate the trend-cycle and seasonal indices. (Experiment with having fixed or changing seasonality.)


```
stl_fixed <- stl(bricksq, s.window = "periodic", robust = TRUE)
autoplot(stl_fixed) + ggtitle("Quarterly Clay Brick Production" , subtitle = "STL Decomposition Fixed S
```

Quarterly Clay Brick Production

STL Decomposition Fixed Seasonality



```
stl_change <- stl(bricksq, s.window = 13, robust = TRUE)
autoplot(stl_change) + ggtitle("Quarterly Clay Brick Production" , subtitle = "STL Decomposition Changing")
```




It is evident that the seasonality changes when it is fixed vs when it is changing.

b. Compute and plot the seasonally adjusted data.

```
brick_stl_adj <- seasadj(stl_fixed)


autoplot(bricksq, series = "Data") +
  autolayer(brick_stl_adj, series = "Seasonally Adjusted") +
  autolayer(trendcycle(stl_fixed), series = "Trend") +
  ggtitle("Quarterly Clay Brick Production", subtitle = "STL Decomposition Fixed Seasonality") +
  ylab("Brick Production") +
  scale_color_manual(values = c("gray", "blue", "red"), breaks = c("Data", "Seasonally Adjusted", "Trend"))
```

c. Use a naïve method to produce forecasts of the seasonally adjusted data.

```
brick_stl_adj %>% naive() %>% autoplot()
```


Forecasts from Naive method



d. Use `stlf()` to reseasonalise the results, giving forecasts for the original data.

```
fcast_brick <- stlf(bricksq, method = "naive")  
autoplot(fcast_brick)
```

Forecasts from STL + Random walk




```
forecast(fcast_brick)
```

##	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## 1994 Q4		465.7338	440.1878	491.2798	426.6645	504.8030
## 1995 Q1		424.6739	388.5464	460.8014	369.4217	479.9261
## 1995 Q2		483.9926	439.7456	528.2395	416.3227	551.6624
## 1995 Q3		494.0000	442.9080	545.0920	415.8616	572.1384
## 1995 Q4		465.7338	408.6112	522.8563	378.3723	553.0952
## 1996 Q1		424.6739	362.0993	487.2485	328.9742	520.3736
## 1996 Q2		483.9926	416.4042	551.5809	380.6251	587.3600
## 1996 Q3		494.0000	421.7450	566.2550	383.4956	604.5044

e. Do the residuals look uncorrelated?

```
checkresiduals(fcast_brick)
```

```
## Warning in checkresiduals(fcast_brick): The fitted degrees of freedom is based
## on the model used for the seasonally adjusted data.
```




```
##
##  Ljung-Box test
##
## data:  Residuals from STL + Random walk
## Q* = 40.829, df = 8, p-value = 2.244e-06
##
## Model df: 0.   Total lags used: 8
```

Looking at the acf and the plot, we can tell the residuals are correlated especially at lag 1, 4, and 8.

f. Repeat with a robust STL decomposition. Does it make much difference?

```
stlf_brick_robust <- stlf(bricksq, method = "naive", robust = TRUE)
autoplot(stlf_brick_robust)
```

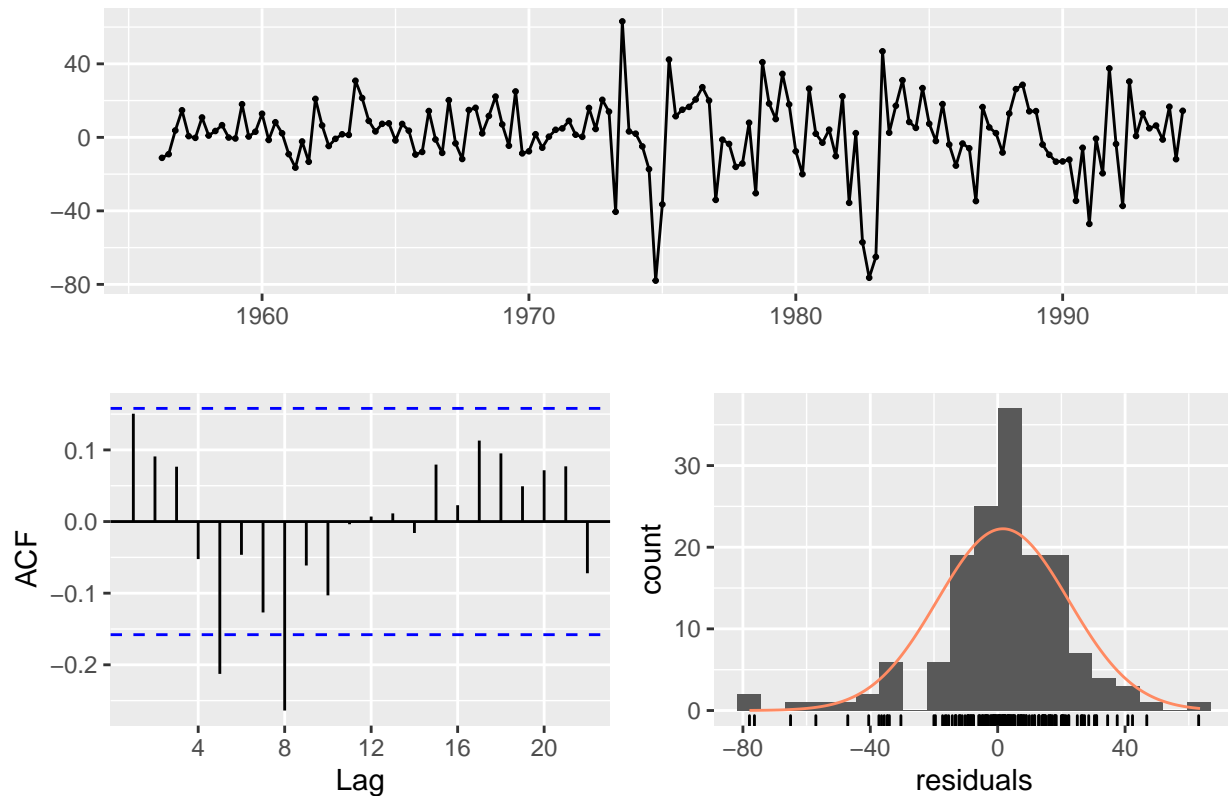
Forecasts from STL + Random walk



```
checkresiduals(stlf_brick_robust)
```

```
## Warning in checkresiduals(stlf_brick_robust): The fitted degrees of freedom is  
## based on the model used for the seasonally adjusted data.
```

Residuals from STL + Random walk



```
##
##  Ljung-Box test
##
## data:  Residuals from STL +  Random walk
## Q* = 27.961, df = 8, p-value = 0.0004817
##
## Model df: 0.   Total lags used: 8
```

the plots are the same, but the correlations are smaller and less.

g. Compare forecasts from `stlf()` with those from `snaive()`, using a test set comprising the last 2 years of data. Which is better?

```
train_brick <- window(bricksq, end= c(1992,3))

test_brick <- window(bricksq, start = c(1992,4))


stlf_train <- stlf(train_brick)

snaive_train <- snaive(train_brick)

autoplot(bricksq, series = "Data") +
  autolayer(snaive_train, PI = FALSE, series = "snaive") + #PI = prediction interval ?
  autolayer(stlf_train, PI = FALSE, series = "stlf") +
  scale_color_manual(values = c("black", "blue", "red"), breaks = c("Data", "snaive", "stlf")) +
  xlim(1990, 1994.75) + ylim(300,600) +
  guides(colour = guide_legend(title = "Data")) +
```

```
ggtitle("Forecast (1992 Q4 - 1994 Q3) with stlf and snaive") +
annotate(
  "rect",
  xmin=1992.75,xmax=1994.5,ymin=-Inf,ymax=Inf,
  fill="#72C6E2",alpha = 0.4)
```

```
## Scale for 'x' is already present. Adding another scale for 'x', which will
## replace the existing scale.
```




It looks like the `stlf()` was better at predicting the forecast. It follows more closely with the original data. `stlf` predicts trend and seasonality better than `snaive`.

Problem 6.7 (Textbook C):


Use `stlf()` to produce forecasts of the writing series with either `method="naive"` or `method="rwdrift"`, whichever is most appropriate. Use the `lambda` argument if you think a Box-Cox transformation is required.

```
autoplot(writing, main= "Writing Series")
```




```
hist(writing)
```

Histogram of writing




```
stlf(writing, robust = TRUE, method = "rwdrift", lambda = BoxCox.lambda(writing)) %>% autoplot()
```




Used a Box-Cox Transformation because the distribution was slightly skewed to the right. I also chose the method `rwdrift` because it performs better than naive for a series that has an overall trend.

Problem 6.8 (Textbook C):


Use `stlf()` to produce forecasts of the fancy series with either `method="naive"` or `method="rwdrift"`, whichever is most appropriate. Use the `lambda` argument if you think a Box-Cox transformation is required.

```
autoplot(fancy, main = "Fancy Series")
```




```
hist(fancy)
```

Histogram of fancy



```
stlf(fancy, robust = TRUE, method = "rwdrift", lambda = BoxCox.lambda(fancy)) %>% autoplot()
```



I used a Box-Cox transformation to make the distribution more evenly distributed and it would perform better forecasts. I chose to use `rwdrift` to perform the forecasts because the original data had an upward trend.